

```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("lakshmi25npathi/imdb-dataset-of-50k-movie-reviews")

print("Path to dataset files:", path)

Using Colab cache for faster access to the 'imdb-dataset-of-50k-movie-reviews' dataset.
Path to dataset files: /kaggle/input/imdb-dataset-of-50k-movie-reviews

!ls /kaggle/input/imdb-dataset-of-50k-movie-reviews

'IMDB Dataset.csv'

import torch
import torch.nn as nn
import pandas as pd
import numpy as np
import re
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from collections import Counter
import matplotlib.pyplot as plt

# -----
# 1 Device
# -----
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

# -----
# 2 Load IMDB dataset
# -----
df = pd.read_csv('/kaggle/input/imdb-dataset-of-50k-movie-reviews/IMDB Dataset.csv')

# Encode labels
le = LabelEncoder()
df['sentiment'] = le.fit_transform(df['sentiment'])

# -----
# 3 Tokenize and Vocabulary
# -----
def tokenize(text):
    return re.findall(r'\b\w+\b', text.lower())

df['tokens'] = df['review'].apply(tokenize)
```

```

all_tokens = [token for tokens in df['tokens'] for token in tokens]
vocab = {word: idx + 2 for idx, (word, _) in
enumerate(Counter(all_tokens).items())}
vocab['<PAD>'] = 0
vocab['<UNK>'] = 1

def encode(tokens):
    return [vocab.get(token, 1) for token in tokens]

df['encoded'] = df['tokens'].apply(encode)

# -----
# 4 Padding
# -----
max_len = 300
def pad_sequence(seq):
    return seq[:max_len] + [0]*(max_len - len(seq)) if len(seq) <
max_len else seq[:max_len]

df['padded'] = df['encoded'].apply(pad_sequence)

# -----
# 5 Train/Test split
# -----
X = np.array(df['padded'].tolist())
y = np.array(df['sentiment'].tolist())

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

X_train = torch.tensor(X_train, dtype=torch.long).to(device)
X_test = torch.tensor(X_test, dtype=torch.long).to(device)
y_train = torch.tensor(y_train, dtype=torch.float32).to(device)
y_test = torch.tensor(y_test, dtype=torch.float32).to(device)

train_data = torch.utils.data.TensorDataset(X_train, y_train)
test_data = torch.utils.data.TensorDataset(X_test, y_test)

train_loader = DataLoader(train_data, batch_size=128, shuffle=True)
test_loader = DataLoader(test_data, batch_size=128)

# -----
# 6 LSTM Model
# -----
class LSTMClassifier(nn.Module):
    def __init__(self, vocab_size, embed_dim=200, hidden_dim=256,
output_dim=1, n_layers=2, bidirectional=True, dropout=0.3):
        super(LSTMClassifier, self).__init__()
        self.embedding = nn.Embedding(vocab_size, embed_dim,
padding_idx=0)

```

```

        self.lstm = nn.LSTM(embed_dim, hidden_dim,
num_layers=n_layers, batch_first=True,
                           bidirectional=bidirectional,
dropout=dropout)
        self.fc = nn.Linear(hidden_dim * 2 if bidirectional else
hidden_dim, output_dim)
        self.sigmoid = nn.Sigmoid()

    def forward(self, x):
        x = self.embedding(x)
        _, (hidden, _) = self.lstm(x)
        if self.lstm.bidirectional:
            hidden = torch.cat((hidden[-2], hidden[-1]), dim=1)
        else:
            hidden = hidden[-1]
        out = self.fc(hidden)
        return self.sigmoid(out)

model = LSTMClassifier(vocab_size=len(vocab)).to(device)

# -----
# 7 Training
# -----
criterion = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer,
'min', patience=2)

epochs = 6
train_losses, test_losses = [], []

for epoch in range(epochs):
    model.train()
    total_loss = 0
    for inputs, labels in train_loader:
        optimizer.zero_grad()
        outputs = model(inputs).squeeze()
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
    train_loss = total_loss / len(train_loader)

    model.eval()
    total_loss = 0
    with torch.no_grad():
        for inputs, labels in test_loader:
            outputs = model(inputs).squeeze()
            loss = criterion(outputs, labels)
            total_loss += loss.item()

```

```

test_loss = total_loss / len(test_loader)

train_losses.append(train_loss)
test_losses.append(test_loss)
scheduler.step(test_loss)

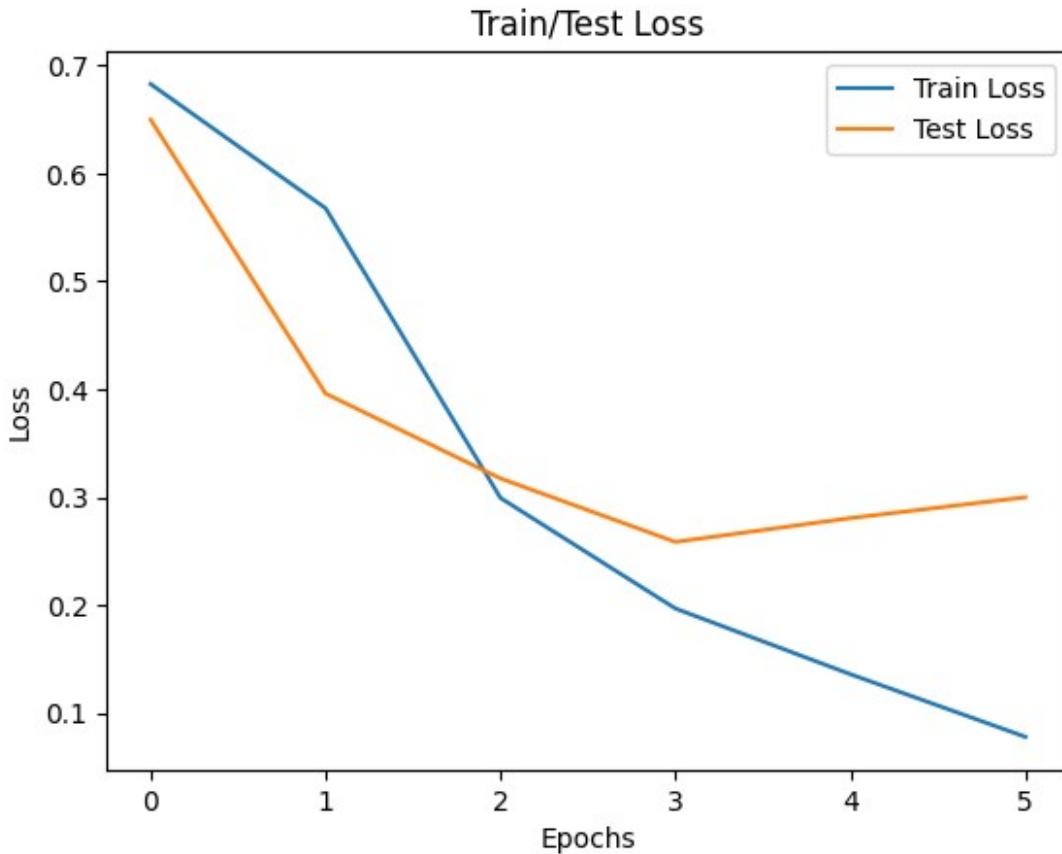
print(f"Epoch {epoch+1}/{epochs} | Train Loss: {train_loss:.4f} | 
Test Loss: {test_loss:.4f}")

# -----
❸ 8 Loss Curves
# -----
plt.plot(train_losses, label='Train Loss')
plt.plot(test_losses, label='Test Loss')
plt.legend()
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Train/Test Loss')
plt.show()

# -----
❹ 9 Evaluate Accuracy
# -----
model.eval()
correct, total = 0, 0
with torch.no_grad():
    for inputs, labels in test_loader:
        outputs = model(inputs).squeeze()
        predicted = (outputs > 0.5).float()
        total += labels.size(0)
        correct += (predicted == labels).sum().item()
accuracy = 100 * correct / total
print(f"Test Accuracy: {accuracy:.2f}%")

Using device: cuda
Epoch 1/6 | Train Loss: 0.6827 | Test Loss: 0.6498
Epoch 2/6 | Train Loss: 0.5677 | Test Loss: 0.3960
Epoch 3/6 | Train Loss: 0.2993 | Test Loss: 0.3174
Epoch 4/6 | Train Loss: 0.1969 | Test Loss: 0.2586
Epoch 5/6 | Train Loss: 0.1360 | Test Loss: 0.2807
Epoch 6/6 | Train Loss: 0.0780 | Test Loss: 0.3000

```



Test Accuracy: 89.77%

```
# □ Prediction
def predict_sentiment(review, model, vocab, max_len, device):
    model.eval()
    tokens = tokenize(review)
    encoded = encode(tokens)
    padded = pad_sequence(encoded)
    input_tensor = torch.tensor(padded,
                                dtype=torch.long).unsqueeze(0).to(device)
    with torch.no_grad():
        output = model(input_tensor).squeeze()
        prediction = (output > 0.5).item()
    return "Positive" if prediction == 1 else "Negative",
    output.item()

# Example prediction
sample_review = "This movie was absolutely amazing! I loved every part
of it."
sentiment, probability = predict_sentiment(sample_review, model,
                                           vocab, max_len, device)
print(f"Review: \'{sample_review}\'")
print(f"Predicted Sentiment: {sentiment} (Probability:
```

```
{probability:.4f})")

sample_review_2 = "This movie was terrible. I hated it."
sentiment_2, probability_2 = predict_sentiment(sample_review_2, model,
vocab, max_len, device)
print(f"Review: \'{sample_review_2}\'")
print(f"Predicted Sentiment: {sentiment_2} (Probability:
{probability_2:.4f})")

Review: "This movie was absolutely amazing! I loved every part of it."
Predicted Sentiment: Positive (Probability: 0.9903)
Review: "This movie was terrible. I hated it."
Predicted Sentiment: Negative (Probability: 0.0013)

model.eval()

LSTMClassifier(
    (embedding): Embedding(101946, 200, padding_idx=0)
    (lstm): LSTM(200, 256, num_layers=2, batch_first=True, dropout=0.3,
bidirectional=True)
    (fc): Linear(in_features=512, out_features=1, bias=True)
    (sigmoid): Sigmoid()
)
```

10/9/25

Exp-8

## Long Short-Term Memory (LSTM) Model

Aim:

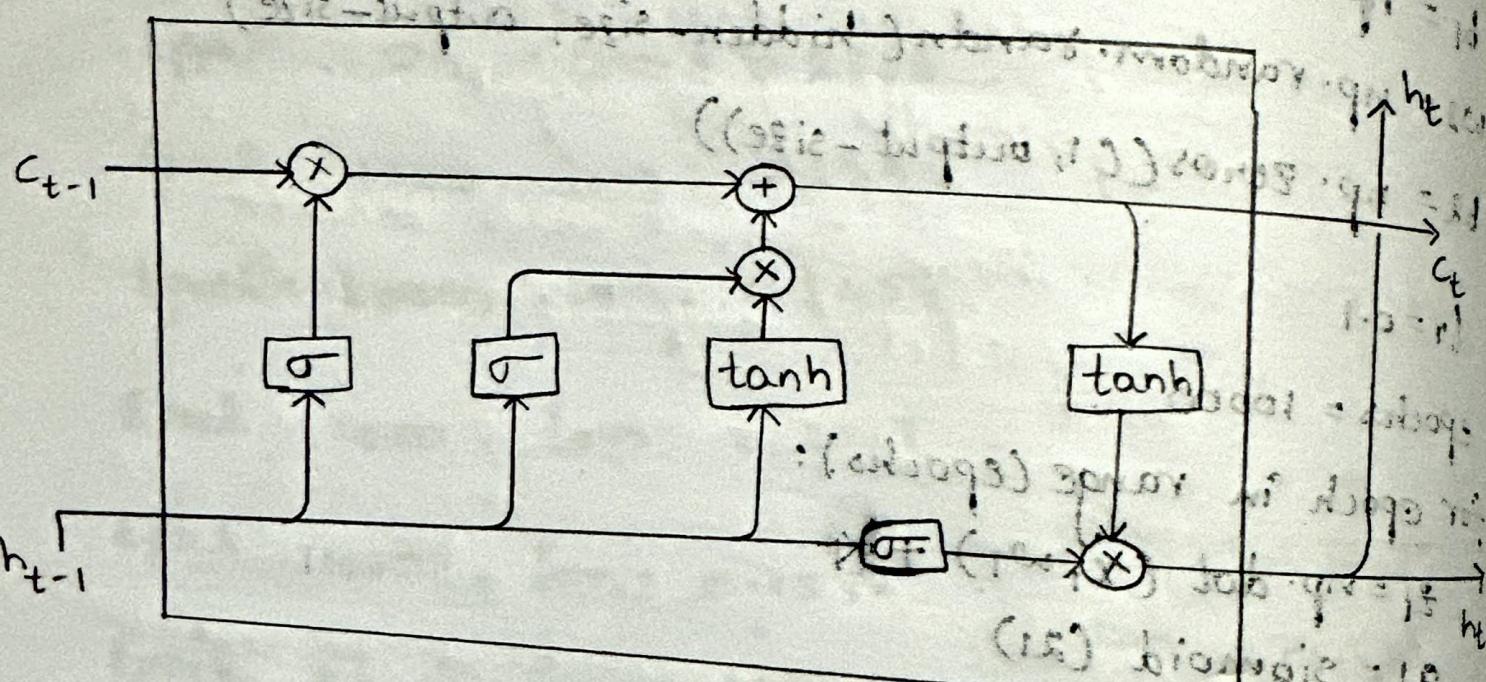
To implement an LSTM model for sequence prediction and observe how it learns temporal dependencies in data.

Description:

LSTMs are a type of RNN designed to capture long-term dependencies in sequential ~~data~~ data. They use memory cells and gates (input, forget and output gates) to regulate information flow, making them effective in tasks like time series prediction, language modeling, and speech

Procedure:

- 1.) Import required libraries and generate a sequence dataset (sine wave).
- 2.) Preprocess the data into input-output pairs for training.
- 3.) Define a LSTM model using keras/Tensorflow.
- 4.) Train the model on the dataset.
- 5.) Evaluate the model and predict future sequence values.



$x_t$

$h_{t-1}$

$c_{t-1}$

$h_t$

$c_t$

$(\sigma)$  biompie

$(\tanh) = 20 - k$

$(\sigma + (\sigma \cdot 10)) \text{ tanh} = 20 - k$

$(1 - (\sigma \cdot 10)) \text{ tanh} = 20 - k$

$(\sigma \cdot 10 + (1 - \sigma \cdot 10)) \text{ tanh} = 20 - k$

$(10 - 10 \cdot \sigma + 10 \cdot \sigma) \text{ tanh} = 20 - k$

$(10 - 10 \cdot \sigma + 10 \cdot \sigma) \text{ tanh} = 20 - k$

$(10 - 10 \cdot \sigma + 10 \cdot \sigma) \text{ tanh} = 20 - k$

code:

```
import numpy as np  
import matplotlib.pyplot as plt  
from tensorflow.keras.models import Sequential  
from tensorflow.keras.layers import LSTM, Dense
```

x = np.linspace(0, 50, 500)

y = np.sin(x)

seq\_length = 20

x, y = [], []

for i in range(len(y) - seq\_length):

x.append(y[i:i + seq\_length])

y.append(y[i + seq\_length])

x = np.array(x)

y = np.array(y)

x = x.reshape(x.shape[0], x.shape[1], 1)

model = Sequential()

model.add(LSTM(50, activation = 'tanh', input\_shape =  
(seq\_length, 1)))

model.add(Dense(1))

model.compile(optimizer = 'adam', loss = 'mse')

history = model.fit(x, y, epochs = 20, verbose = 0)

pred = model.predict(x, verbose = 0)

plt.plot(y, label = "Actual")

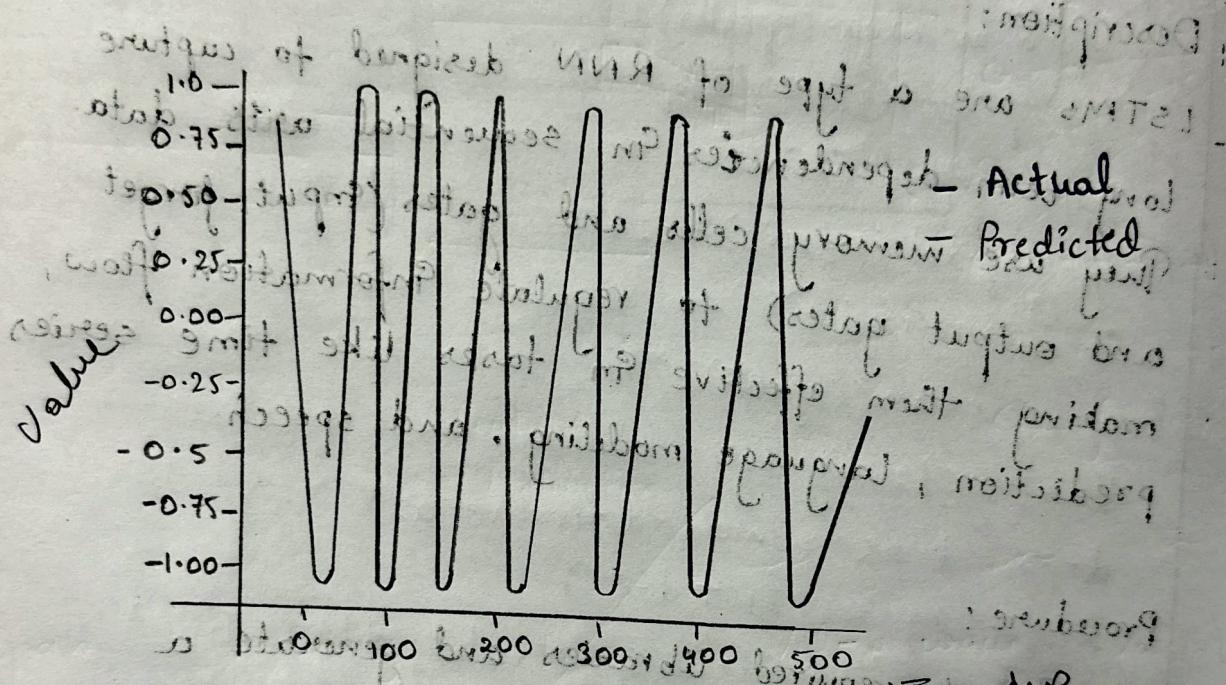
plt.plot(pred, label = "Predicted")

plt.legend()

plt.show()

Output:

- 1.) Training loss (MSFD) decreases during epochs.
- Epoch 1/20 - Loss: ~0.05
- Epoch 20/20 - Loss: ~0.0012



- 2.) Graph output - The plot shows the actual sine wave and the predicted sine wave by LSTM.
- The black curve (actual sine wave)
- The grey curve (predicted sine wave by LSTM)

Result:

- 1.) The LSTM model successfully learned the sine wave patterns with a final loss close to zero.
- 2.) The predicted wave overlapped with the actual curve, proving the model's ability to capture sequential dependencies.